

A CNN Chess Engine

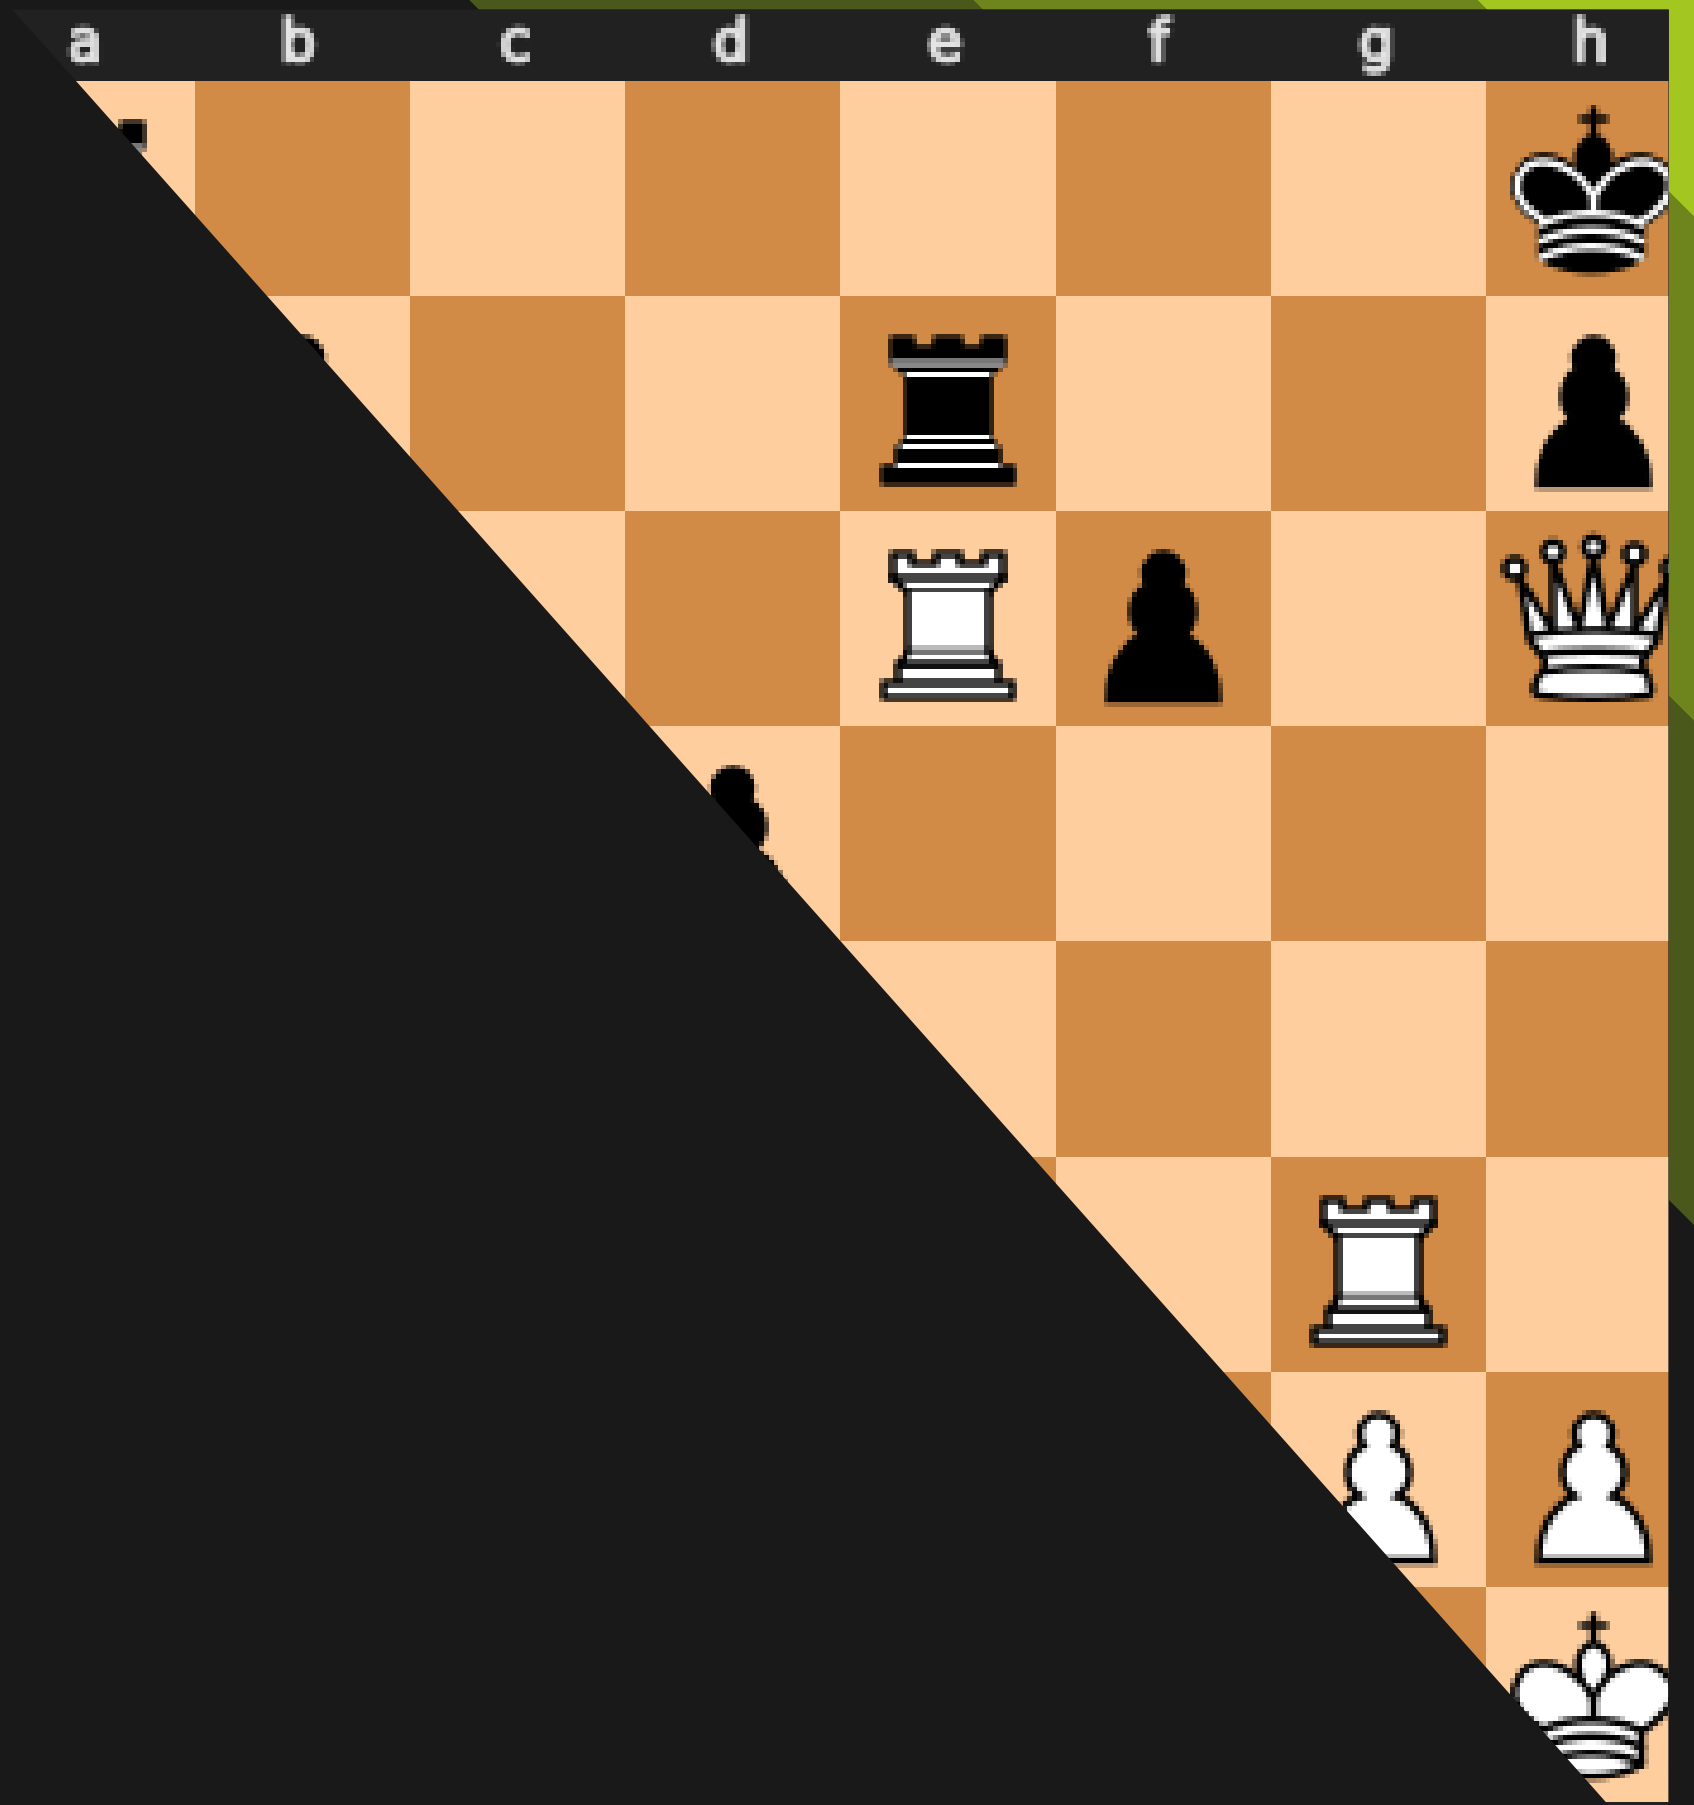
Nathaniel Martin

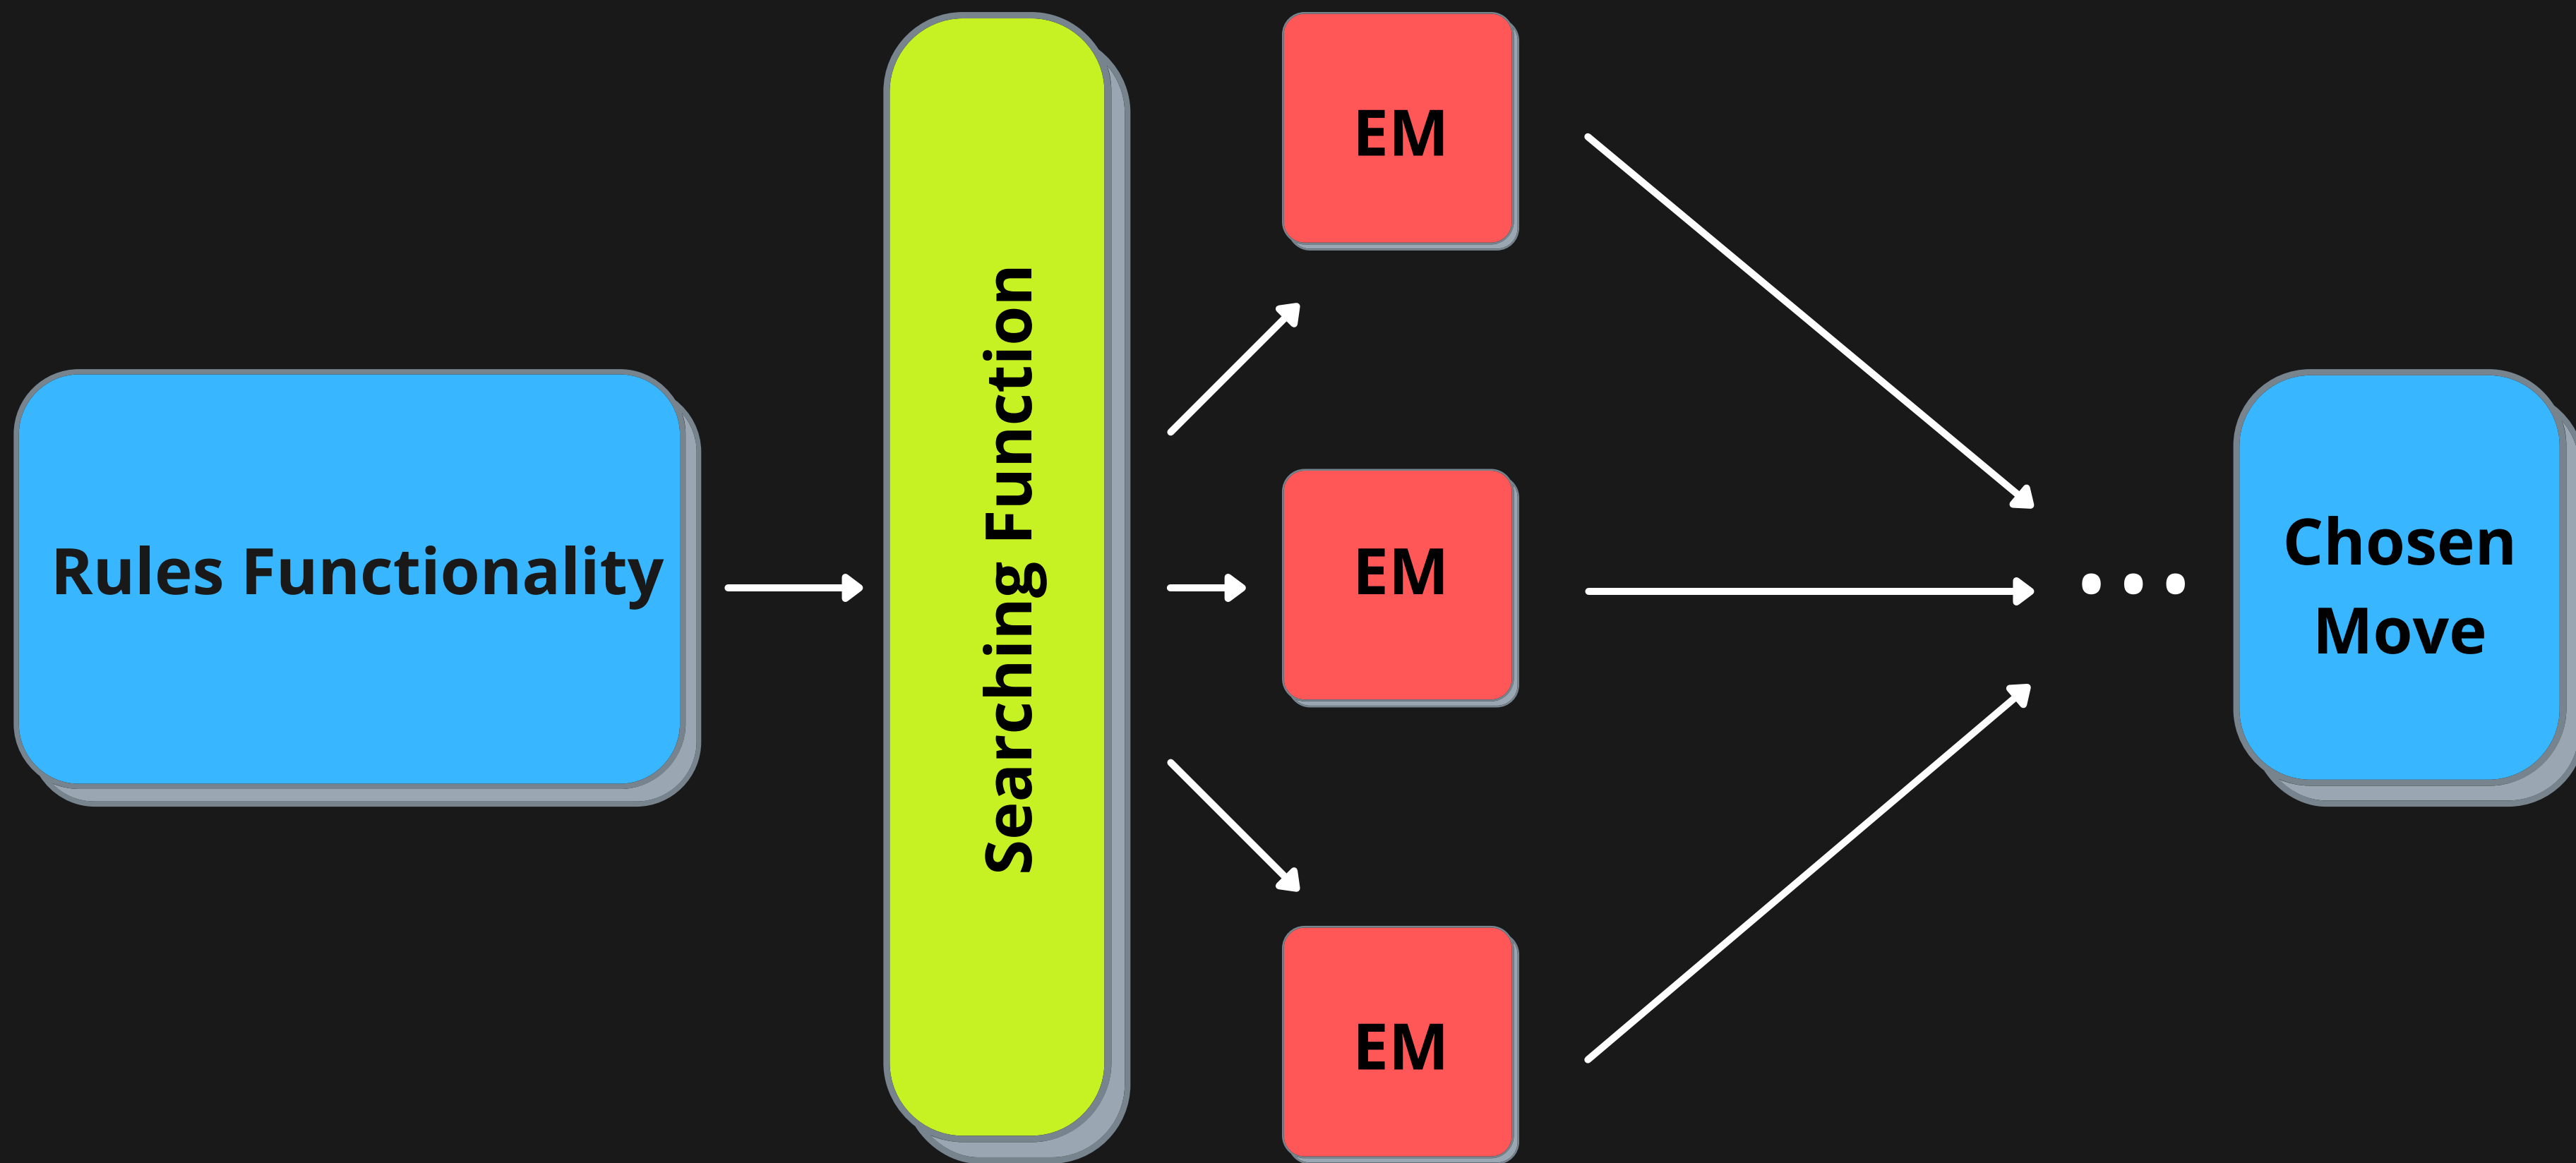


Agenda

- ▶ Engine Architecture
- ▶ Data Preparation
- ▶ Model Construction
- ▶ Performance

Engine Architecture





Data Preparation

Lichess Puzzle Data



~ 2.2 million puzzles

All taken from actual games played.



A position has ~ 32 possible moves

This makes the data very imbalanced



"Only Move" Puzzles

This lets us label moves easily.



FEN positions and UCI moves


Each entry has a starting FEN and move list. The puzzle position is actually the starting FEN advanced by the first move.

Puzzle Pieces

f2g3

UCI Moves

The first move is a prep move to get to the actual puzzle position. The second move is our target.



r6k/pp2r2p/4Rp1Q/3p4/8/1N1P2R1/PqP2bPP/7K




r6k/pp2r2p/4Rp1Q/3p4/8/1N1P2b1/PqP3PP/7K



What the Network Gets

- Each "slice" of the board has one piece type.
- This allows each slice to be represented by an 8 x 8 grid of ones and zeros.



Model Construction



Convolutional Neural Network

Typically used in Image Classification


Relatively Small

Just under 55,000 parameters

Binary Classification

Thanks to our puzzle data

Performance



Typical Metrics



Accuracy

Evaluation metrics need to give as accurate as possible scores to the search function.



Naive Imbalanced: ~96%

Notably this is slightly *worse* than guessing that every move is bad, given 32 moves on average.



Longer Training, and Imbalance change: ~88%

With 1 in 5 positions being good moves, this is doing better than a dumb model.

To see more...

Tensorboard:

Play the bot on Lichess:

<https://lichess.org/@/PuzzledBot>





Thank you!

Linkedin: www.linkedin.com/in/nathaniel-martin-73b037227

Github: <https://github.com/UpGoerFive>